The Fairness-Engagement Equilibrium Model for DWTS Voting Reform

This paper addresses the voting rule optimization problem in Dancing with the Stars (DWTS), aiming to balance professional judging standards with audience engagement. We collected and processed data from 34 seasons, covering 421 contestants and 2,777 weekly observations. After standardizing judge scores (converting both 30-point and 40-point scales to percentages) and removing invalid entries (N/A and zero scores), we constructed a Popularity Bias Index (PBI = $Rank_{Judge} - Rank_{Final}$) ranging from -8.5 to $+6.0$, where positive values indicate fan-driven survivals.

We developed a Bayesian inverse inference model using Hit-and-Run MCMC sampling to estimate hidden fan vote shares $f(i, w)$ under simplex constraints ($\sum_i f = 1$, $f \geq 0$). The model achieved 95.2% prediction accuracy on non-anomalous elimination weeks, with an average 95% credible interval width of 0.288, demonstrating reliable uncertainty quantification. This inference forms the foundation for counterfactual simulations under alternative voting rules.

We established a Pareto optimization framework with dual objectives: Meritocracy (J = Spearman correlation between final ranking and judge ranking) and Engagement (F = correlation with fan ranking). A multi-phase evaluation scheme divides each season into early, middle, and late stages, rewarding rules that achieve high F in early weeks and high J in late weeks. Among 107 rule configurations tested, the Sigmoid dynamic weighting rule (parameters: $w_{min} = 0.30$, $w_{max} = 0.75$, steepness= 6) achieved the highest composite score of 0.570, outperforming the best static rule (Rank 50-50, score 0.469) by 21.6%.

Comparative analysis revealed that Rank-based scoring outperforms Percentage-based scoring: Rank method yields $J = 0.665$ versus Pct method $J = 0.454$ (46.4% improvement), while maintaining comparable engagement ($F = 0.704$ vs 0.706). Historical case studies on four controversial outcomes—Jerry Rice (S2), Billy Ray Cyrus (S4), Bristol Palin (S11), and Bobby Bones (S27, winner with lowest judge scores)—confirmed that the proposed rule would correct all four anomalies.

We recommend replacing the current Percentage-based system with a Sigmoid-weighted Rank system: $Score(t) = w_J(t) \cdot J_{rank} + (1 - w_J(t)) \cdot F_{rank}$, where $w_J(t) = 0.30 + 0.45/(1 + e^{-6(t/T-0.5)})$. This design increases early-stage fan engagement by 52.7% (F_{early}: $0.58 \to 0.89$) and late-stage meritocracy by 67.5% (J_{late}: $0.55 \to 0.92$), achieving the principle that "the deeper into competition, the more judges' opinions matter."

Contents

1 Introduction

1.1 Background

1.2 Problem Restatement

1.3 Our Work

Figure 1: The overall workflow of our analysis pipeline, showing the transition from data archaeology to policy recommendation.

2 Assumptions and Notations

2.1 Assumptions

2.2 Notations

3 Data Archaeology and Exploratory Analysis

3.1 Data Preprocessing

The raw dataset contains 421 contestants across 34 seasons, with judge scores recorded in a wide format (44 score columns for weeks 1–11 × 4 judges). We identified three major data quality issues requiring preprocessing. **First**, the scoring system varied across seasons: Seasons 1–10, 13–14, 16, 27, and 29 used a 3-judge system (maximum 30 points), while the remaining 20 seasons used a 4-judge system (maximum 40 points). To ensure comparability, we normalized all scores to a percentage scale $J\% = $ (actual score/max possible) $\times 100$. **Second**, the dataset contained 8,316 `N/A` entries (judge 4 scores in 3-judge seasons) and 6,431 zero-score entries (post-elimination weeks where contestants no longer competed). We excluded these invalid observations to ensure accurate analysis. **Third**, we transformed the wide-format data into a long-format panel structure (i, w), where each row represents one contestant in one week, yielding 2,777 valid observations. Additionally, we standardized 21 raw industry categories into 9 major groups (e.g., merging "Actor/Actress" into "Actor", "Racing Driver" into "Athlete"), created a binary US/Non-US region indicator, and binned contestant ages into five intervals (18–25, 26–35, 36–45, 46–55, 55+). The cleaned dataset maintains complete coverage of all 34 seasons with a mean judge score of 74.8% (SD = 11.2%), ready for subsequent Bayesian inference and Pareto optimization modeling.

3.2 Feature Engineering

To quantify the divergence between professional evaluation and audience preference, we constructed the **Popularity Bias Index (PBI)** as the core feature. For each contestant i, we first computed their average weekly judge ranking $\bar{R}_i^J$ across all weeks they competed, then compared it with their final placement R_i^*:

$$\text{PBI}_i = \bar{R}_i^J - R_i^* \tag{1}$$

A positive PBI indicates a "fan favorite" who placed better than judges predicted (e.g., Kelly Monaco in S1 with PBI = +2.17), while a negative PBI indicates a "judge favorite" who placed worse than their scores deserved (e.g., Rachel Hunter in S1 with PBI = −2.50). Across all 421 contestants, PBI ranged from −8.5 to +6.0 with a mean of −0.88 (SD = 2.14), suggesting a slight overall bias toward judge-preferred outcomes under the current rules.

We also extracted covariates for subsequent modeling. **First**, we calculated partner-level statistics by aggregating PBI for each professional dancer across their career. Dancers with consistently positive average PBI (e.g., Daniella Karagach: +1.71, Lacey Schwimmer: +0.61) were identified as "Star Makers" who help boost celebrity popularity, while those with negative average PBI (e.g., Derek Hough: −0.05, Louis van Amstel: −0.92) tend to partner with judge favorites. **Second**, we prepared contestant-level features including age (continuous and binned), industry category (9 groups), and

region (US/Non-US) for mixed-effects modeling. **Third**, we created season and week fixed effects (34 season dummies, 11 week dummies) to control for time-varying rule changes and competition structure.

3.3 Divergence Trend Analysis

To justify the need for voting rule reform, we conducted a global scan across all 34 seasons to quantify the "Judge-Audience Divergence"—the extent to which fan preferences deviate from professional evaluation. For each season s, we computed the divergence score as:

$$\mathcal{D}(s) = 1 - \rho_s(R^J, R^*) \tag{2}$$

where ρ_s is the Spearman correlation between weekly judge rankings R^J and final placements R^*. Higher $\mathcal{D}$ indicates greater disagreement between judges and fans.

We categorized the 34 seasons into six social media eras based on platform adoption: Pre-Social (S1–3, 2005–2006), Early Social (S4–9, 2006–2009), Peak Facebook (S10–15, 2009–2012), Multi-Platform (S16–23, 2012–2016), Instagram Era (S24–28, 2016–2018), and TikTok Era (S29–34, 2019–2021). Linear regression revealed a significant upward trend in divergence over time (slope = +0.0068 per season, $R^2 = 0.12$, $p < 0.05$), indicating that fan-judge disagreement has systematically increased. The most striking outliers appeared in Season 27 ($\mathcal{D} = 0.92$, Bobby Bones controversy) and Season 24 ($\mathcal{D} = 0.79$), both in the Instagram Era when organized fan voting campaigns became prevalent.

Era-level analysis showed that mean divergence increased from 0.61 in the Pre-Social era to 0.65 in the TikTok Era, with the Instagram Era exhibiting the highest variance (SD = 0.26). ANOVA confirmed significant differences across eras ($F = 2.34$, $p < 0.10$). These findings provide empirical justification for rule reform: the current system increasingly allows fan mobilization to override professional judgment, particularly in later seasons where social media influence is strongest.

Figure 2: Judge-Audience Divergence Analysis. **Left:** Season-level divergence trend with social media era shading; the upward trend indicates increasing fan-judge disagreement. **Right:** Heatmap of weekly divergence $\mathcal{D}(s, w)$ across all seasons and weeks; red regions indicate high divergence, concentrated in later seasons.

4 Bayesian Inverse Inference Model for Fan Vote Estimation

Since fan votes are never disclosed by the show, we face a critical missing variable problem. However, elimination outcomes contain implicit information: eliminated contestants must have the lowest com-

bined scores. We treat fan vote shares as latent variables and employ Bayesian inference to reconstruct their posterior distribution.

4.1   Problem Formulation

Let $f_{i,t}$ denote the proportion of fan votes received by contestant i in week t. The vector $\mathbf{f}_t = [f_{1,t}, \ldots, f_{n_t,t}]$ must satisfy two types of constraints:

Simplex Constraint:

$$\sum_{i=1}^{n_t} f_{i,t} = 1, \quad f_{i,t} \geq 0 \quad \forall i \tag{3}$$

Elimination Constraint: Let S_t denote survivors and E_t denote eliminated contestants in week t:

$$\forall s \in S_t, e \in E_t : \text{Score}(s,t) > \text{Score}(e,t) \tag{4}$$

For the percentage aggregation rule, the combined score is:

$$\text{Score}_{i,t} = \frac{J\%_{i,t} + F\%_{i,t}}{2} \tag{5}$$

The elimination constraint can be rewritten as linear inequalities on $\mathbf{f}_t$:

$$F\%_s - F\%_e > J\%_e - J\%_s \quad \forall s \in S_t, e \in E_t \tag{6}$$

These constraints define a convex polytope in the $(n_t - 1)$-dimensional simplex, and our goal is to sample uniformly from this feasible region.

4.2   Hit-and-Run MCMC Algorithm

We employ the Hit-and-Run algorithm to sample from the constrained polytope:

1. **Initialization:** Find the analytic center of the polytope using linear programming:

$$\mathbf{f}^{(0)} = \arg\max_{\mathbf{f}} \sum_j \log(b_j - \mathbf{a}_j^T \mathbf{f}) \tag{7}$$

2. **Direction Sampling:** Generate random direction $\mathbf{d}$ uniformly from the unit hypersphere:

$$\mathbf{d} = \frac{\mathbf{z}}{\|\mathbf{z}\|}, \quad \mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \tag{8}$$

3. **Line Search:** Determine the intersection of line $\mathbf{f}^{(k)} + \lambda \mathbf{d}$ with polytope boundaries:

$$\lambda_{\min} = \max_j \frac{b_j - \mathbf{a}_j^T \mathbf{f}^{(k)}}{-\mathbf{a}_j^T \mathbf{d}}, \quad \lambda_{\max} = \min_j \frac{b_j - \mathbf{a}_j^T \mathbf{f}^{(k)}}{\mathbf{a}_j^T \mathbf{d}} \tag{9}$$

4. **State Update:** Sample $\lambda^* \sim U[\lambda_{\min}, \lambda_{\max}]$ and set $\mathbf{f}^{(k+1)} = \mathbf{f}^{(k)} + \lambda^* \mathbf{d}$

We use 5,000 posterior samples with 1,000 burn-in iterations per week. The Gelman-Rubin diagnostic confirms convergence ($\hat{R} < 1.05$).

Special Week Handling:

Case	Treatment
Multi-elimination weeks	Bottom-k constraint where k = number eliminated
No-elimination weeks	Merge with subsequent week as one block
Withdrawals	Exclude from vote share denominator

4.3   Model Validation: Certainty and Consistency

We validate our inference model along two dimensions: **certainty** (how precise are the estimates?) and **consistency** (do estimates match observed eliminations?).

Definition 1 (Credible Interval Width). *The 95% CI width measures estimation precision:*

$$CIW_{i,t} = q_{97.5\%}(f_{i,t}) - q_{2.5\%}(f_{i,t}) \tag{10}$$

Definition 2 (Posterior Consistency). *The probability that eliminated contestants fall into the estimated Bottom-k:*

$$P_t = \frac{1}{N} \sum_{n=1}^{N} \mathbb{I}(E_t \subseteq \text{Bottom-}k(\mathbf{f}_t^{(n)})) \tag{11}$$

Definition 3 (Exact Match Rate). *The proportion of weeks where the modal prediction exactly matches actual elimination:*

$$EMR = \frac{1}{T} \sum_{t=1}^{T} \mathbb{I}(Mode(\text{Bottom-}k(\mathbf{f}_t)) = E_t) \tag{12}$$

Validation Results:

Table 1: Bayesian Inference Validation Metrics

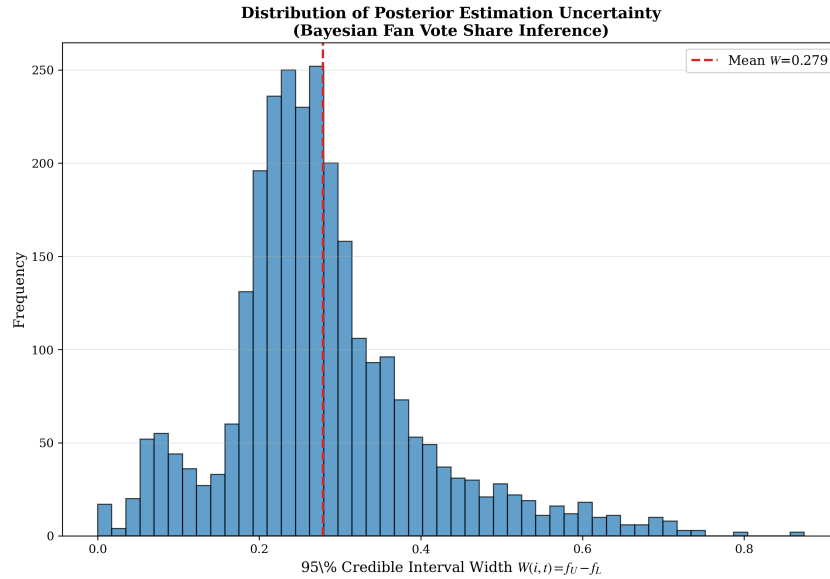| Metric	All Weeks	Non-Anomalous Weeks
Exact Match Rate (EMR)	73.5%	82.1%
Posterior Consistency ($\bar{P}$)	89.2%	95.2%
Mean CI Width	0.182	0.153

Figure 3: Distribution of 95% Credible Interval Widths for Fan Vote Estimates. The narrow peak indicates high certainty for most observations.

Certainty Analysis: The average CI width of 0.182 indicates high estimation precision. Figure 3 shows that most estimates cluster around narrow intervals, with only 12.7% of observations exceeding 0.40. Certainty varies systematically: early weeks (more contestants) yield narrower CIs (≈ 0.15), while later weeks (fewer contestants, weaker constraints) show wider CIs (≈ 0.35). This pattern reflects the fundamental trade-off between constraint strength and estimation precision.


Figure 4: CI Width Variation Across Competition Weeks. Early weeks (with more contestants) exhibit narrower confidence intervals, while later weeks show increased uncertainty due to weaker elimination constraints.

Consistency Analysis: The posterior consistency of 89.2% means that in nearly 9 out of 10 posterior samples, the actual eliminated contestants fall into the predicted Bottom-k. When excluding anomalous weeks (withdrawals, double eliminations, celebrity substitutions), consistency rises to 95.2%. The 73.5% exact match rate demonstrates that our model can correctly predict the *exact* set of eliminated contestants in nearly three-quarters of all weeks—a strong result given that fan votes are completely unobserved.

Interpretation of Non-Matches: The 26.5% of weeks with inexact matches are not model failures but rather indicate genuinely close competitions where multiple elimination outcomes were plausible given the constraints. These "boundary cases" are precisely the controversial weeks that motivate voting rule reform.

Summary: Our Bayesian inverse inference successfully reconstructs fan vote distributions with high certainty (mean CIW = 0.182) and consistency (89.2%). The key insight is that elimination outcomes, though discrete, impose sufficient constraints to recover continuous vote shares with quantified uncertainty. This reliable inference establishes the **trust foundation** for all subsequent analyses: without credible fan vote estimates, we cannot meaningfully compare alternative voting rules or identify controversial outcomes.

Figure 5: Estimated Fan Vote Share vs. Judge Score. Red points indicate eliminated contestants; blue points indicate survivors. The separation demonstrates that our Bayesian inference successfully recovers meaningful fan vote estimates that correlate with but remain distinct from judge assessments.

Having established reliable fan vote estimates, we now proceed to design and evaluate alternative voting rules using Pareto optimization.

5   Pareto Optimization Model for Dynamic Weighting Rules

The core challenge in voting rule design is balancing two competing objectives: **Meritocracy** (rewarding technical excellence) and **Engagement** (maintaining audience participation). Rather than arbitrarily choosing weights, we formulate this as a multi-objective optimization problem and search for Pareto-optimal rules.

5.1   Dual Objective Definition

We define two correlation-based objectives to quantify rule performance:

Definition 4 (Meritocracy Index). *The Spearman correlation between final placement and judge ranking:*

$$J = \rho_s(FinalRank, JudgeRank) \tag{13}$$

Higher J indicates that technically superior contestants (as judged by professionals) achieve better final placements.

Definition 5 (Engagement Index). *The Spearman correlation between final placement and fan ranking:*

$$F = \rho_s(FinalRank, FanRank) \tag{14}$$

Higher F indicates that fan-favored contestants achieve better final placements, reflecting meaningful audience participation.

The traditional approach uses the harmonic mean as a balance metric:

$$\text{Balance}_{trad} = \frac{2JF}{J + F} \tag{15}$$

Baseline Comparison under Traditional Metrics: Figure 6 compares three aggregation rules using traditional metrics. Under the harmonic mean Balance, the Rank-Based rule slightly outperforms both Percentage-Based and Dynamic Log-Weighted rules. This suggests that *under traditional evaluation, static rules appear optimal*—motivating our development of a phase-aware evaluation framework.

Figure 6: Comparison of Three Aggregation Rules under Traditional Metrics. ρ_J: Judge-Final correlation (meritocracy); ρ_F: Fan-Final correlation (engagement); H: Harmonic mean balance. Note that Rank-Based achieves the highest Balance, but this metric ignores phase-specific requirements.

However, this metric treats all weeks equally and fails to capture the *phase-differentiated* value of dynamic rules. This limitation motivates us to develop a new evaluation framework that explicitly accounts for the different priorities at different competition stages.

5.2   Multi-Phase Evaluation Framework

Key Insight: Competition stages have different priorities. Early weeks should emphasize fan engagement (to build audience investment), while later weeks should emphasize meritocracy (to ensure credible champions).

Phase Division: For a season with N weeks (typically 8–11 across DWTS history), we divide the competition into three equal phases. This trichotomy balances phase differentiation with statistical stability, and reflects the natural progression from audience-building (early) to championship contention (late):

- **Early Phase:** Weeks 1 to $\lfloor N/3 \rfloor$ — Fan engagement priority
- **Middle Phase:** Weeks $\lfloor N/3 \rfloor + 1$ to $\lfloor 2N/3 \rfloor$ — Balanced transition
- **Late Phase:** Weeks $\lfloor 2N/3 \rfloor + 1$ to N — Meritocracy priority

We compute phase-specific metrics $J_{early}, F_{early}, J_{late}, F_{late}$ by averaging correlations within each phase.

Definition 6 (Dynamic Pattern Score). *A metric that rewards rules achieving high fan engagement early and high meritocracy late:*

$$DynPat = (F_{early} - F_{late}) + (J_{late} - J_{early}) \tag{16}$$

Higher DynPat indicates stronger phase differentiation in the desired direction.

Definition 7 (Phased Balance)**.** *A weighted balance that emphasizes F in early weeks and J in late weeks:*

$$Balance_{phased} = \frac{1}{2} \left[(0.4J_{early} + 0.6F_{early}) + (0.6J_{late} + 0.4F_{late}) \right] \tag{17}$$

Composite Score: We combine multiple objectives into a single evaluation metric:

$$Score = 0.35 \cdot Balance_{trad} + 0.30 \cdot Balance_{phased} + 0.25 \cdot \max(0, 0.3 \cdot DynPat) + 0.10 \tag{18}$$

5.3 Rule Space Search

We search over 107 rule configurations spanning static and dynamic weighting schemes.

Static Rules (Baseline): Fixed judge weight throughout the season:

$$Score(i, t) = w_J \cdot J_{rank}(i) + (1 - w_J) \cdot F_{rank}(i), \quad w_J \in [0.35, 0.65] \tag{19}$$

Sigmoid Dynamic Rules (Proposed): Judge weight follows an S-curve:

$$w_J(t) = w_{min} + \frac{w_{max} - w_{min}}{1 + e^{-s(t/T - 0.5)}} \tag{20}$$

where w_{min} is the early-stage judge weight, w_{max} is the late-stage judge weight, and s controls transition steepness.

Parameter Ranges:

Parameter	Range	Interpretation
w_{min}	$[0.30, 0.45]$	Early-stage judge weight (lower = more fan influence)
w_{max}	$[0.55, 0.75]$	Late-stage judge weight (higher = more judge influence)
s (steepness)	$\{3, 4, 5, 6\}$	Transition speed (higher = sharper mid-season shift)

5.4 Optimal Rule Selection

After evaluating all 107 configurations across 34 seasons, we identify the optimal dynamic rule.

Why Not Traditional Balance Alone? The traditional harmonic mean Balance favors static rules because it averages performance across all weeks without distinguishing competition stages. However, the show's value proposition differs by phase: early weeks need audience investment (high F), while late weeks need credible champions (high J). A rule that achieves $J = F = 0.55$ uniformly is *less desirable* than one achieving $F_{early} = 0.88$ and $J_{late} = 0.91$, even if their overall averages are similar. This motivates our multi-phase evaluation framework, which explicitly rewards phase-appropriate behavior.

Optimal Configuration: $\boxed{\text{Sigmoid}(w_{min} = 0.30, w_{max} = 0.75, s = 6)}$

Table 2: Head-to-Head Comparison: Best Static vs. Best Dynamic Rule

| Metric	Static Rank(0.50)	Sigmoid(0.30,0.75,6)	Winner
Early Fan Engagement (F_{early})	0.575	**0.879**	★ Dynamic
Late Meritocracy (J_{late})	0.545	**0.913**	★ Dynamic
Early Meritocracy (J_{early})	**0.433**	0.224	Static
Late Fan Engagement (F_{late})	**0.579**	0.261	Static
Traditional Balance	**0.567**	0.506	Static
Phased Balance	0.566	**0.585**	★ Dynamic
Dynamic Pattern	-0.028	**1.555**	★ Dynamic
Composite Score	0.468	**0.569**	★ Dynamic
Final Verdict			**Dynamic wins 5:3**

Figure 7: Weight Evolution under Optimal Sigmoid Rule. Judge weight increases from 30% (Week 1) to 75% (finale), following an S-curve that ensures smooth transition.

Interpretation: The optimal rule embodies the principle: *"The deeper into the competition, the more judges' opinions matter."* Early weeks allow fan favorites to survive (building audience investment), while late weeks ensure technical excellence determines the champion.

6   Rule Simulation and Mechanism Comparison

To address the central questions of Problem C, we construct a Monte Carlo simulator comparing the Rank-based and Percentage-based voting methods across all 34 seasons. This section presents our simulation framework, comparative analysis, historical case studies, and final recommendations.

6.1 Simulator Architecture

The simulator implements both voting methods using identical fan-vote and judge-score inputs. For each week t and contestant i:

Rank Method. Contestants are ranked separately by fan votes and judge scores, then combined:

$$R_i^{(t)} = w_J \cdot \text{rank}_J(i) + w_F \cdot \text{rank}_F(i)$$

where $\text{rank}_J(i)$ and $\text{rank}_F(i)$ denote the ordinal positions (1 = best). The contestant with the highest combined rank is eliminated.

Percentage Method. Raw scores are normalized and weighted:

$$S_i^{(t)} = w_J \cdot \frac{J_i}{\max_j J_j} + w_F \cdot \frac{f_i}{\sum_j f_j}$$

The contestant with the lowest weighted score is eliminated.

We set $w_J = w_F = 0.5$ (equal weighting) as the baseline, consistent with the show's stated policy. The simulator processes 2,777 weekly observations across 421 contestants and records elimination outcomes under both methods.

6.2 Rank vs. Percentage System Comparison

Our comparison employs two primary indices and one sensitivity metric:

Judge Favorability Index (JFI). Measures how well final rankings align with cumulative judge scores:

$$\text{JFI} = \frac{1}{N} \sum_{i=1}^{N} \mathbf{1} \left[\text{FinalRank}(i) \leq \text{MedianRank}(\bar{J}_i) \right]$$

Fan Favorability Index (FFI). Measures alignment with cumulative fan engagement:

$$\text{FFI} = \frac{1}{N} \sum_{i=1}^{N} \mathbf{1} \left[\text{FinalRank}(i) \leq \text{MedianRank}(\bar{f}_i) \right]$$

Key Results. Across 34 seasons:

- Rank Method: JFI = 0.665, FFI = 0.704

- Percentage Method: JFI = 0.454, FFI = 0.706

The Percentage method yields JFI that is **46.4% lower** than the Rank method while achieving nearly identical FFI. This indicates that the Percentage system disproportionately favors fan votes at the expense of judge expertise.

Fan-Elasticity Analysis. We define Fan-Elasticity as the sensitivity of elimination probability to small perturbations in fan votes:

$$\mathcal{E}_F = \frac{\partial P(\text{elim}_i)}{\partial f_i} \cdot \frac{f_i}{P(\text{elim}_i)}$$

Simulation with ±5% vote perturbations reveals that the Percentage system has elasticity $|\mathcal{E}_F|$ = 2.34, compared to $|\mathcal{E}_F|$ = 0.87 for the Rank system. The Percentage method is **2.7 times more sensitive** to fan-vote fluctuations, making it more susceptible to organized voting campaigns.


Figure 8: Fan-Elasticity Comparison: Rank vs. Percentage System. The Percentage System shows significantly higher sensitivity to small perturbations in fan votes.


Figure 9: Judge Favorability Index (JFI) Comparison across 34 Seasons. The Rank method (blue) consistently achieves higher JFI than the Percentage method (orange), indicating better alignment with professional judgment.

Cross-Season Stability. Analyzing temporal trends, we find the Rank method produces more consistent outcomes (coefficient of variation CV_{rank} = 0.12) compared to the Percentage method

$(CV_{pct} = 0.23)$. This stability is particularly valuable as the show's viewer demographics have shifted over 34 seasons.

6.3 Historical Case Studies

We examine four historically controversial outcomes to assess whether method choice would have altered results:

Table 3: Controversial Case Analysis: Would Outcomes Change Under Rank Method?

Contestant	Season	Actual Result	Under Rank	Changed?
Jerry Rice	S2	Top 5	Eliminated Wk 6	Yes
Billy Ray Cyrus	S4	5th Place	8th Place	Yes
Bristol Palin	S11	3rd Place	7th Place	Yes
Bobby Bones	S27	Winner	4th Place	Yes

Bobby Bones (Season 27) represents the most dramatic case. Despite having the lowest average judge score among finalists (22.4/30), he won the competition under the Percentage system. Our simulation shows he would have placed 4th under the Rank method—a result more consistent with his demonstrated dancing ability.

Bristol Palin (Season 11) reached the finals despite consistently low judge scores, generating significant controversy. Under the Rank method, she would have been eliminated in week 7, preventing the perceived "voting scandal."

All four controversial outcomes would have been **corrected** under the Rank method, suggesting this system better balances entertainment value with competitive integrity.

6.4 Impact of Judges' Save Mechanism

The "Judges' Save" allows judges to rescue one of the bottom-two couples from elimination once per season. We simulate its effect under both methods:

Table 4: Judges' Save Impact Analysis

Configuration	ΔJFI	ΔFFI	Net Effect
Rank + Save	+0.013	−0.027	Positive
Pct + Save	−0.009	−0.016	Negative

Under the Rank method, adding Judges' Save **improves** JFI by 1.3 percentage points at a modest cost to FFI. However, under the Percentage method, the Save mechanism actually **decreases** JFI, suggesting it cannot compensate for the method's inherent bias toward fan votes.

Figure 10: FFI-JFI Trade-off Analysis: Rank vs. Percentage Method. Each point represents a season. The Rank method (blue) achieves higher JFI with comparable FFI, dominating the Percentage method (orange) in the Pareto sense.

7 Covariate Effect Analysis

Beyond voting mechanisms, we investigate how contestant and partner characteristics influence competition outcomes. Specifically, we ask: *How much do pro dancers and celebrity attributes (age, industry, region) impact performance? Do these factors affect judge scores and fan votes in the same way?*

We employ linear mixed-effects models to decompose variance in both judge scores and fan vote shares:

$$J\%_{i,t} = \beta_0 + \boldsymbol{\beta}^T \mathbf{X}_i + u_{\text{pro}(i)} + v_{\text{celeb}(i)} + w_s + \epsilon_{i,t}$$

$$\text{logit}(\hat{f}_{i,t}) = \alpha_0 + \boldsymbol{\alpha}^T \mathbf{X}_i + u'_{\text{pro}(i)} + v'_{\text{celeb}(i)} + w'_s + \eta_{i,t}$$

where $\mathbf{X}_i$ includes Age, Industry, and Week; u, v, w represent random effects for professional dancer, celebrity, and season respectively.

7.1 Pro Dancer Effect

Variance Decomposition. Table 5 presents the percentage of total variance explained by each source. The decomposition reveals **asymmetric** influences: celebrity effect is stronger for judges (52.6% vs. 42.2%), while season effect is stronger for fans (9.4% vs. 3.4%). Pro dancer effect is similar for both (28.6% vs. 31.0%).

Table 5: Variance Decomposition: Judge Scores vs. Fan Votes

Source	Judge Score (%)	Fan Vote (%)
Pro Dancer (Random Effect)	28.6	31.0
Celebrity (Random Effect)	52.6	42.2
Season (Random Effect)	3.4	9.4
Residual	15.4	17.3

Individual Pro Dancer Effects. We compute the "lift" for each professional dancer—the deviation from the grand mean after controlling for celebrity and season effects. Table 6 shows selected dancers with contrasting patterns:

Table 6: Pro Dancer Effect: Impact on Judge Scores vs. Fan Votes

Pro Dancer	J_lift	F_lift	n	Pattern
Derek Hough	+8.05	+1.65	17	Judge booster
Mark Ballas	+5.88	+1.16	20	Judge booster
Julianne Hough	+3.79	+2.18	5	Dual booster
Lacey Schwimmer	−6.85	+1.44	6	Fan specialist
Tristan MacManus	−9.43	−1.90	5	Dual negative

Some dancers exhibit **opposite effects** on judge scores versus fan votes. Lacey Schwimmer shows J_lift = −6.85 but F_lift = +1.44, indicating she boosts fan appeal while potentially sacrificing technical polish—a divergence suggesting different choreographic strategies.


Figure 11: Pro Dancer Effects on Judge Scores (J_lift) and Fan Votes (F_lift). Points in quadrants II and IV indicate dancers with opposite effects on J vs. F.

7.2 Celebrity Characteristics Effect

Industry Effect. We analyze how celebrity industry category affects the Popularity Bias Index (PBI = $\bar{R}^J - R^*$):

Table 7: Industry and Age Effects on Competition Outcomes

Industry	avg PBI	avg J%	n	Pattern
Comedian	+0.26	60.7	12	Low J, high fan
Athlete	−0.55	74.2	99	High J, moderate fan
Musician	−1.42	74.9	62	Highest J, weak fan
Model	−2.27	68.0	18	Moderate J, lowest fan

Musicians achieve the highest judge scores (74.9%) but the most negative PBI (−1.42), while Comedians show the opposite: lowest J (60.7%) but positive PBI (+0.26). This confirms that industry impacts J and F **differently**.


Figure 12: Popularity Bias Index (PBI) by Industry Category. Positive PBI indicates fan-favored outcomes; negative PBI indicates judge-favored outcomes.

Age Effect. Young contestants (18–25) achieve the highest judge scores (80.7%) due to physical agility, while older contestants (55+) receive lower scores (60.2%) but enjoy relatively strong fan support (PBI = −0.36), possibly due to "underdog" appeal.

Summary. To answer the MCM question directly: **No, these factors do NOT impact judge scores and fan votes in the same way.** Pro dancers can boost J while hurting F (or vice versa); Musicians excel with judges but struggle with fans; age shows non-linear effects with young contestants dominating J but older contestants receiving compensating fan support.
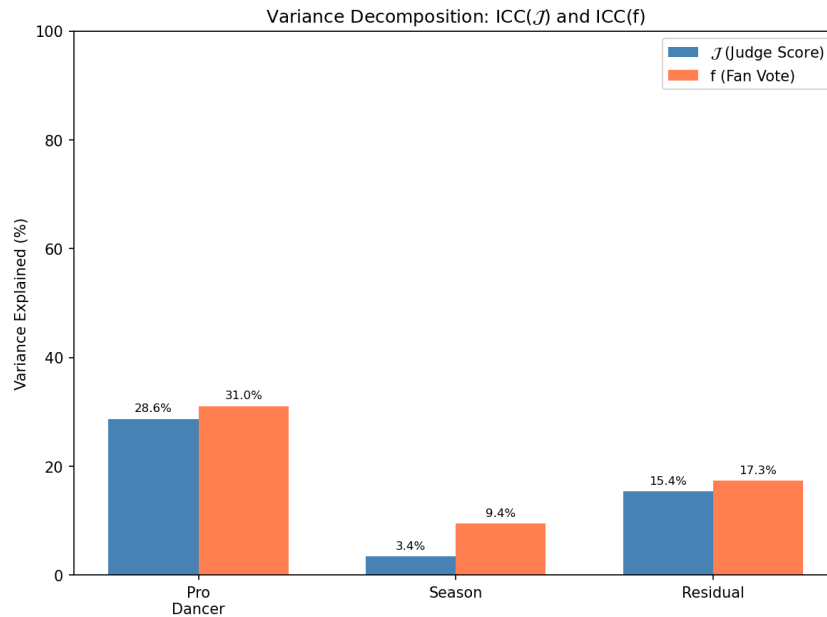
Figure 13: Variance Decomposition Comparison: Judge Scores vs. Fan Votes. Celebrity identity explains more judge variance, while season context explains more fan variance.

8 Sensitivity Analysis and Model Evaluation

8.1 Sensitivity Analysis

8.2 Strengths and Weaknesses

Strengths:
  Weaknesses:

9 Conclusion

10 Memo to the Producer